

Introduction to Statistical Learning with Applications

CM3: Cross-validation, model selection,
and bias-variance

Pedro L. C. Rodrigues

IN OUR PREVIOUS EPISODES...

Our current workflow

- 1 We want to estimate the values of Y based on predictors X_1, \dots, X_p
- 2 We are given a set of N examples of y_i with corresponding x_{i1}, \dots, x_{ip}
- 3 We estimate parameters $\hat{\beta}$ that minimize $\frac{1}{N} \sum_{i=1}^N (y_i - \beta^T x_i)^2$
- 4 We do statistical inference on the values of the $\hat{\beta}_j$

$$\downarrow x_i = \begin{bmatrix} x_{i1} \\ \vdots \\ x_{ip} \end{bmatrix}$$

IN OUR PREVIOUS EPISODES...

We do **statistical inference** on the values of $\hat{\beta}_j$

1

We are given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$

3

We assume the noise is **Gaussian** IID

$$\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$


2

We assume that each data point is

$$y_i = \beta_0 + \sum_{k=1}^p \beta_k x_{ik} + \varepsilon_i$$

$$\hat{\beta} = \begin{bmatrix} \hat{\beta}_0 \\ \hat{\beta}_1 \\ \vdots \\ \hat{\beta}_p \end{bmatrix}$$

$$\longrightarrow \hat{\beta}_i \sim \mathcal{N}(\beta_i, \sigma_i^2)$$



IN OUR PREVIOUS EPISODES...

We do **statistical inference** on the values of $\hat{\beta}_j$


Note that even if we had data generated with $\beta_i = 0$ our estimate is a RV

$$\hat{\beta}_i \sim \mathcal{N}(0, \sigma^2) \longrightarrow \text{Very likely that } \hat{\beta}_i \neq 0$$

Therefore, we would like to check whether the number that we get for this estimate is one that could **indeed come from such distribution**

If it is a **number too far off from the confidence interval**, then it might in fact mean that the distribution with zero-mean is not a good one to describe the data

This is the intuition behind the statistical tests!



IN OUR PREVIOUS EPISODES...

However, in practice we're often more interested in the **predictions!**

Is there a natural way of describing the statistics of the predicted values?

Remember that if we assume a linear model for the data, then


$$\mathbb{E}[Y \mid X = x] = m(x) = \beta_0 + \sum_{k=1}^p \beta_k x_k = x^\top \beta$$

Using our limited dataset, we can build

$$\hat{m}(x) = x^\top \hat{\beta} \sim \mathcal{N}(x^\top \beta, \Sigma_m) \text{ with } \Sigma_m = x^\top \Sigma_\beta x$$


IN OUR PREVIOUS EPISODES...

But what if we can't assume a linear model? Nor Gaussian?




IN OUR PREVIOUS EPISODES...

But what if we can't assume a linear model? Nor Gaussian?



IN OUR PREVIOUS EPISODES...

But what if we can't assume a linear model? Nor Gaussian?



IN OUR PREVIOUS EPISODES...

But what if we can't assume a linear model? Nor Gaussian?

arXiv > cs > arXiv:2107.07511

Computer Science > Machine Learning

[Submitted on 15 Jul 2021 (v1), last revised 7 Dec 2022 (this version, v6)]

A Gentle Introduction to Conformal Prediction and Distribution-Free Uncertainty Quantification

Anastasios N. Angelopoulos, Stephen Bates

 awesome-conformal-prediction Public

 Watch 11  Fork 44  Star 526

 main  1 Branch  1 Tag


 Go to file

t

Add file

 Code

About

 valeman Update README.md

329ea75 · yesterday  1,099 Commits

 documents

Add files via upload

2 years ago


A professionally curated list of awesome Conformal Prediction videos, tutorials, books, papers, PhD and MSc theses, articles and open-source libraries.

IN OUR PREVIOUS EPISODES...

Our current workflow


- 1 We want to estimate the values of Y based on predictors X_1, \dots, X_p
- 2 We are given a set of N examples of y_i with corresponding x_{i1}, \dots, x_{ip}
- 3 We estimate parameters $\hat{\beta}$ that minimize $\frac{1}{N} \sum_{i=1}^N (y_i - \beta^T x_i)^2$
- 4 We do statistical inference on the values of the $\hat{\beta}_j$
- 5 How can we **assess the quality** of a model?

$$\downarrow x_i = \begin{bmatrix} x_{i1} \\ \vdots \\ x_{ip} \end{bmatrix}$$



- What makes a model **good**?
- Estimating the quality of a model
- Comparing and selecting models

What makes a model **good**?



Suppose our data is generated as


$$Y = \beta_0 + \sum_{i=1}^d \beta_i X^i + \varepsilon$$

we can run multiple linear regression with

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^d \end{bmatrix}$$

What's the best choice for d ?

What makes a model **good**?



Suppose our data is generated as


$$Y = \beta_0 + \sum_{i=1}^d \beta_i X^i + \varepsilon$$

we can run multiple linear regression with

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^d \end{bmatrix}$$

What's the best choice for d ?

What makes a model **good**?



Suppose our data is generated as

$$Y = \beta_0 + \sum_{i=1}^d \beta_i X^i + \varepsilon$$

we can run multiple linear regression with

$$\mathbf{X} = \begin{bmatrix} 1 & x_1 & x_1^2 & \dots & x_1^d \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_N & x_N^2 & \dots & x_N^d \end{bmatrix}$$

What's the best choice for d ?

What makes a model **good**? – The generalization error

Remember that our goal in regression is to minimize the generalization error:



$$\mathcal{L}(r) = \mathbb{E}_{Y,X} \left[\left(Y - r(X) \right)^2 \right]$$

Intuitively, we can expect this error to be decomposed into a few informative parts:

- The **irreducible error**: can we ever predict Y from X with zero prediction error? Probably not, since we always assume having some observation error ε in the data model.
- The **estimation bias**: when estimating the conditional expectation, we always have to choose a family of approximators, which may not be sufficiently flexible.
- The **estimation variance**: for families of approximators that are too flexible, it might happen that for every small change in the observed data points, the estimates change a lot.

What makes a model **good**?

Estimation variance. The same model fit on two slightly different datasets.



What makes a model **good**?



What makes a model **good**? – The bias-variance tradeoff


Remember that the data model is assumed to be $Y = r(X) + \varepsilon$ with $\text{Var}(\varepsilon) = \sigma^2$

We use a training dataset $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$ to estimate $\hat{r}_{\mathcal{D}}$

For new observation x we write the bias-variance decomposition conditioned on x as

$$\begin{aligned}\mathcal{L}(\hat{r}_{\mathcal{D}}(x)) &= \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \mid X = x \right] \\&= \mathbb{E}_{Y,X} \left[\left(Y - r(X) + r(X) - \hat{r}_{\mathcal{D}}(X) \right)^2 \mid X = x \right] \\&= \mathbb{E}_{Y,X} \left[\left(Y - r(X) \right)^2 \mid X = x \right] + \mathbb{E}_{Y,X} \left[\left(r(X) - \hat{r}_{\mathcal{D}}(X) \right)^2 \mid X = x \right] \\&= \underbrace{\mathbb{E}_{Y,X} \left[\left(Y - r(X) \right)^2 \mid X = x \right]}_{\text{Irreducible error } \sigma^2} + \underbrace{\left(\mathbb{E}_{Y,X} \left[r(X) - \hat{r}_{\mathcal{D}}(X) \mid X = x \right] \right)^2}_{\text{Squared estimation bias}} + \underbrace{\text{Var}(\hat{r}_{\mathcal{D}}(X) \mid X = x)}_{\text{Estimation variance}}\end{aligned}$$

- What makes a model **good**?



- Estimating the quality of a model

- Comparing and selecting models

Estimating the quality of a model

How can we estimate the generalization error of an approximator model in practice?

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right]$$



We don't
know the $p(Y, X)$

Estimating the quality of a model

How can we estimate the generalization error of an approximator in practice?

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right] \approx \frac{1}{M} \sum_{i=1}^M \left(y_i - \hat{r}_{\mathcal{D}}(x_i) \right)^2 = L(\hat{r}_{\mathcal{D}}, \mathcal{X}) \quad \text{with } \mathcal{X} = \left\{ (y_i, x_i) \right\}_{i=1}^M$$



We don't
know the $p(Y, X)$

So we approximate it
with M data points

Note that for $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^N$ we have $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$

We want to minimize one
quantity but can only estimate
a proxy of it


One-slide reminder of the law of large numbers (LLN)

Whenever we want to approximate an expectation, we can use the LLN

$$\mathbb{E}_X [f(X)] = \int f(x)p(x)dx \approx \frac{1}{N} \sum_{n=1}^N f(x_i)$$

with N **independent** and identically distributed samples $x_i \sim p(x)$

This result will be useful many times in our class and is ubiquitous in machine learning



This is also used in computer graphics, statistical physics, econometry, etc.

Estimating the quality of a model

How can we estimate the generalization error of an approximator in practice?

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{Y,X} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right] \approx \frac{1}{M} \sum_{i=1}^M \left(y_i - \hat{r}_{\mathcal{D}}(x_i) \right)^2 = L(\hat{r}_{\mathcal{D}}, \mathcal{X}) \quad \text{with } \mathcal{X} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=M}$$



We don't know the $p(Y, X)$

So we approximate it with M data points

Note that for $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$ we have $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$

We want to minimize one quantity but can only estimate a proxy of it

Question:

Can we say that $\mathcal{L}(\hat{r}_{\mathcal{D}}) \approx L(\hat{r}_{\mathcal{D}}, \mathcal{D})$?

ABSOLUTELY NOT!



In fact $\mathcal{L}(\hat{r}_{\mathcal{D}}) \geq L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ i.e. the true generalization error is larger than the estimated one

The optimism of the training error

Different ways of seeing the problem:

- Intuitively, since $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$ then $L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ will always look small...

The optimism of the training error


Different ways of seeing the problem:

- Intuitively, since $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$ then $L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ will always look small...
- Consider our example with polynomials
 - Increasing the degree improves training error
 - But the testing error grows!

$L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ is the **training** error

$L(\hat{r}_{\mathcal{D}}, \mathcal{X})$ with $\mathcal{X} \neq \mathcal{D}$ is the **testing** error

Error on unseen data



The optimism of the training error


Different ways of seeing the problem:

- Intuitively, since $\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} L(r, \mathcal{D})$ then $L(\hat{r}_{\mathcal{D}}, \mathcal{D})$ will always look small...
- Consider our example with polynomials
 - Increasing the degree improves training error
 - But the testing error grows!
- Mathematical illustration on linear regression

$$\mathcal{D} = \left\{ (x_i, y_i) \right\}_{i=1}^{i=N} \quad \mathcal{D}' = \left\{ (x_i, y'_i) \right\}_{i=1}^{i=N} \quad Y = \beta_0 + \sum_{k=1}^p \beta_k X_k + \varepsilon$$

(same predictors but different observations because of noise)

$$\mathbb{E} \left[L(\hat{r}_{\mathcal{D}}, \mathcal{D}') \right] = \mathbb{E} \left[L(\hat{r}_{\mathcal{D}}, \mathcal{D}) \right] + \frac{2}{N} \sigma^2 (p + 1)$$



The optimism of the training error

1

We are given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$

The optimism of the training error

1

We are given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$

2

We train regressor on \mathcal{D} that minimizes

$$\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left(y_i - r(x_i) \right)^2$$

The optimism of the training error

1

We are given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$

3

We want to estimate the generalization error of $\hat{r}_{\mathcal{D}}$ which is defined as per

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{(X,Y)} \left[(Y - \hat{r}_{\mathcal{D}}(X))^2 \right]$$

2

We train regressor on \mathcal{D} that minimizes

$$\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left(y_i - r(x_i) \right)^2$$

The optimism of the training error

1

We are given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$

3

We want to estimate the generalization error of $\hat{r}_{\mathcal{D}}$ which is defined as per

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{(X,Y)} \left[(Y - \hat{r}_{\mathcal{D}}(X))^2 \right]$$

2

We train regressor on \mathcal{D} that minimizes

$$\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left(y_i - r(x_i) \right)^2$$

4

But we can only obtain an approximation

$$L(\hat{r}_{\mathcal{D}}, \mathcal{X}) = \frac{1}{M} \sum_{i=1}^M \left(y_i - \hat{r}_{\mathcal{D}}(x_i) \right)^2$$

such that $\mathcal{X} = \{x_i, y_i\}_{i=1}^M$
 $(x_i, y_i) \sim p(x, y)$

The optimism of the training error

1

We are given a dataset $\mathcal{D} = \{x_i, y_i\}_{i=1}^N$

3

We want to estimate the generalization error of $\hat{r}_{\mathcal{D}}$ which is defined as per

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{(X,Y)} \left[(Y - \hat{r}_{\mathcal{D}}(X))^2 \right]$$

!

Beware the optimism of training error!

$$L(\hat{r}_{\mathcal{D}}, \mathcal{D}) \leq L(\hat{r}_{\mathcal{D}}, \mathcal{X})$$

This was the criterion
used during the training

2

We train regressor on \mathcal{D} that minimizes

$$\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{H}} \frac{1}{N} \sum_{i=1}^N \left(y_i - r(x_i) \right)^2$$

4

But we can only obtain an approximation

$$L(\hat{r}_{\mathcal{D}}, \mathcal{X}) = \frac{1}{M} \sum_{i=1}^M \left(y_i - \hat{r}_{\mathcal{D}}(x_i) \right)^2$$

such that $\mathcal{X} = \{x_i, y_i\}_{i=1}^M$
 $(x_i, y_i) \sim p(x, y)$


The optimism of the training error

Conclusion:

- We can only know if a model is good or not if we **correctly evaluate** its performance
- The objective function that we minimize when training a model (**training error**) is never the same as the one we are actually interested in minimizing (**test error**)

$$\hat{r}_{\mathcal{D}} = \operatorname{argmin}_{r \in \mathcal{F}} \frac{1}{M} \sum_{i=1}^M \left(y_i - r(x_i) \right)^2$$

$$\mathcal{L}(\hat{r}_{\mathcal{D}}) = \mathbb{E}_{(X,Y)} \left[\left(Y - \hat{r}_{\mathcal{D}}(X) \right)^2 \right]$$



Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

- **Strategy 1 : Single hold-out test point**

We fit a model on the first N-1 training samples, calling it $\hat{r}^{(-N)}$

Treat the last sample (X_N, Y_N) as a test sample and estimate generalization error as

$$\mathcal{L}(\hat{r}) \approx \left(y_N - \hat{r}^{(-N)}(x_N) \right)^2$$

The estimator is easy to calculate but we can expect a rather large variance.

Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

- **Strategy 2 : Leave one out cross-validation (LOOCV)**

We fit N models on N-1 training samples, holding out x_i at each time to get $\hat{r}^{(-i)}$

The test error is approximate as

$$\mathcal{L}(\hat{r}) \approx \frac{1}{N} \sum_{i=1}^N \left(y_i - \hat{r}^{(-i)}(x_i) \right)^2$$

The variance decreases, but the computational burden is much higher

Cross-validation

We are given a dataset with N datapoints $\mathcal{D} = \left\{ (y_i, x_i) \right\}_{i=1}^{i=N}$

How can we estimate the **generalization error** of a model we train on this dataset?

Remember: We should evaluate the model on samples that were not used for training

○ Strategy 3 : K-Fold cross-validation

Split the training dataset randomly into K folds so to have $\mathcal{D}_1 \cup \dots \cup \mathcal{D}_K = \mathcal{D}$

For $k=1, \dots, K$ fit a model $\hat{r}^{(-k)}$ on a training set except that excludes \mathcal{D}_k

$$\mathcal{L}(\hat{r}) \approx \frac{1}{K} \sum_{k=1}^K \left(\underbrace{\frac{1}{N_k} \sum_{i \in \mathcal{D}_k} \left(y_i - \hat{r}^{(-k)}(x_i) \right)^2}_{\text{Estimate test error for each fold}} \right)$$

Take average
over folds



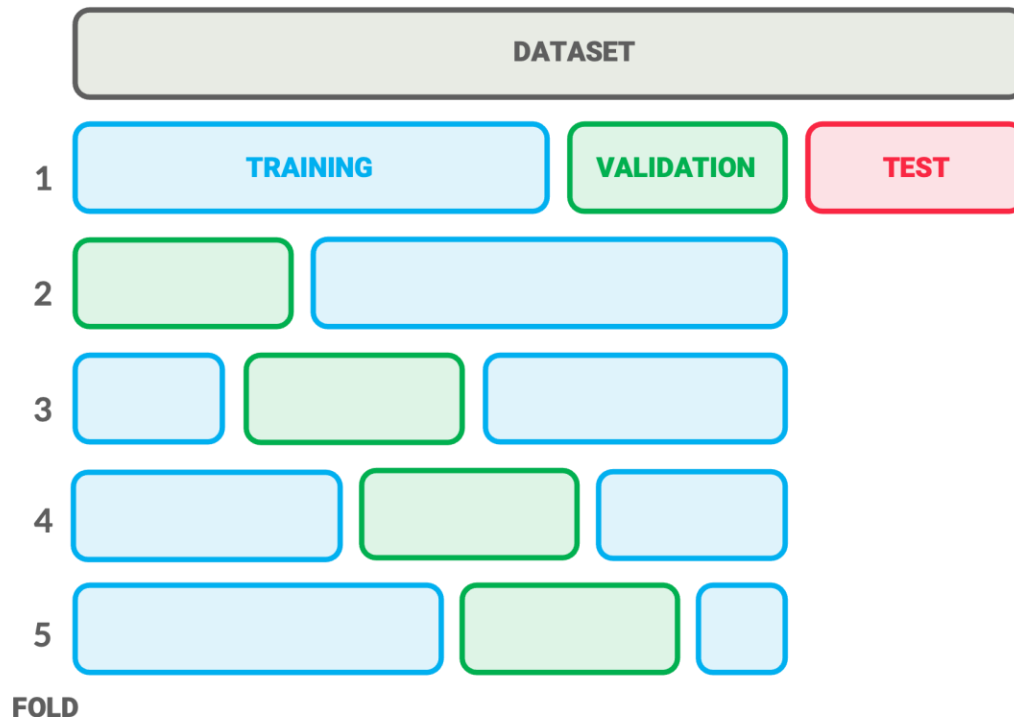
Estimate test error for each fold

Cross-validation

Strategy 3 : K-Fold cross-validation

Note that the errors calculated in each fold are not IID random variables, since the models share some of their data points.

↳ **Question:** What does this imply?



Using 

```
from sklearn.model_selection import KFold
from sklearn.datasets import fetch_california_housing
from sklearn.linear_model import LinearRegression
kf = KFold(n_splits=5)
dataset = fetch_california_housing()
lm = LinearRegression()
X = dataset.data
y = dataset.target
scores = []
for idx_train, idx_test in kf.split(X):
    X_train, y_train = X[idx_train], y[idx_train]
    X_test, y_test = X[idx_test], y[idx_test]
    lm.fit(X_train, y_train)
    scores.append(lm.score(X_test, y_test))
print(np.mean(scores))
```

See the documentation for more details

Cross-validation

But there are **several** other strategies!

<u>GroupKFold</u>	K-fold iterator variant with non-overlapping groups.
<u>GroupShuffleSplit</u>	Shuffle-Group(s)-Out cross-validation iterator.
<u>KFold</u>	K-Fold cross-validator.

<u>ShuffleSplit</u>	Random permutation cross-validator.
<u>StratifiedGroupKFold</u>	Stratified K-Fold iterator variant with non-overlapping groups.
<u>StratifiedKFold</u>	Stratified K-Fold cross-validator.
<u>StratifiedShuffleSplit</u>	Stratified ShuffleSplit cross-validator.
<u>TimeSeriesSplit</u>	Time Series cross-validator.

and more...

Cross-validation

(1) Example with the **categorical variables** from last week

```
from sklearn.model_selection import KFold, ShuffleSplit
from sklearn.linear_model import LinearRegression

regressor = LinearRegression()
cv = KFold(n_splits=4)
results_cv = cross_validate(
    regressor, X, y, cv=cv, scoring='neg_mean_squared_error')
print(f'MSE: {-results_cv["test_score"].mean():.2f}')

regressor = LinearRegression()
cv = ShuffleSplit(n_splits=4)
results_cv = cross_validate(
    regressor, X, y, cv=cv, scoring='neg_mean_squared_error')
print(f'MSE: {-results_cv["test_score"].mean():.2f}')
```

17.50

4.60

In [85]: df


Out[85]:

	life	rpm	brand
0	18.73	610	A
1	14.52	950	A
2	17.43	720	A
3	14.54	840	A
4	13.44	980	A
5	24.39	530	A
6	13.34	680	A
7	22.71	540	A
8	12.68	890	A
9	19.32	730	A
10	30.16	670	B
11	27.09	770	B
12	25.40	880	B
13	26.05	1000	B
14	33.49	760	B
15	35.62	590	B
16	26.07	910	B
17	36.78	650	B
18	34.95	810	B
19	43.67	500	B

What is going on? 😄

Cross-validation

(2) Example with **financial** time series: predict CVX's quotes based on other quotes



Cross-validation

(2) Example with **financial** time series: predict CVX's quotes based on other quotes

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import ShuffleSplit, cross_validate
X = quotes.drop(columns=["CVX"])
y = quotes["CVX"]
regressor = GradientBoostingRegressor()
cv = ShuffleSplit(n_splits=10)
results_cv = cross_validate(regressor, X, y, cv=cv)
print(f'Mean R2: {results_cv["test_score"].mean():.2f}')
```

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y}_i)^2}$$

0.95

This looks like an almost perfect prediction! 🤖 Does it sound right to you?

Cross-validation

(2) Example with **financial** time series: predict CVX's quotes based on other quotes

```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import ShuffleSplit, cross_validate
X = quotes.drop(columns=["CVX"])
y = quotes["CVX"]
regressor = GradientBoostingRegressor()
cv = ShuffleSplit(n_splits=10)
results_cv = cross_validate(regressor, X, y, cv=cv)
print(f'Mean R2: {results_cv["test_score"].mean():.2f}')
```

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

0.95

This looks like an almost perfect prediction! 🤖 Does it sound right to you?

```
from sklearn.model_selection import TimeSeriesSplit
cv = TimeSeriesSplit(n_splits=10)
results_cv = cross_validate(regressor, X, y, cv=cv)
print(f'Mean R2: {results_cv["test_score"].mean():.2f}')
```

-3.10

Disappointing, but closer to reality...

Cross-validation


(3) Choosing the best order for the polynomial approximating a dataset

```
1  import matplotlib.pyplot as plt
2  import numpy as np
3  from sklearn.preprocessing import PolynomialFeatures
4  from sklearn.pipeline import make_pipeline
5  from sklearn.linear_model import LinearRegression
6  from sklearn.model_selection import GridSearchCV, ShuffleSplit
7
8  np.random.seed(6)
9
10 # generate dataset
11 N = 100
12 x = np.linspace(-5, +5, N)
13 y = x - x**2 + x**3 + 20 * np.random.randn(N)
14 X = x.reshape(-1, 1)
15
16 # instantiate pipeline
17 poly = PolynomialFeatures(include_bias=False)
18 lr = LinearRegression()
19 pipe = make_pipeline(poly, lr)
20
21 # instantiate grid search cv
22 degrees_array = np.arange(1, 10+1)
23 parameters = {'polynomialfeatures__degree':degrees_array}
24 cv = ShuffleSplit(n_splits=10)
25 est = GridSearchCV(estimator=pipe, param_grid=parameters, cv=cv,
26                    scoring='neg_mean_squared_error', return_train_score=True)
27
28 # get cv results
29 est.fit(X, y)
30 cv_test_scores = -est.cv_results_['mean_test_score']
31 cv_train_scores = -est.cv_results_['mean_train_score']
32
33 # plot results
34 fig, ax = plt.subplots(figsize=(6, 5.4))
35 ax.plot(degrees_array, cv_test_scores, c='C0', lw=3, label='test')
36 ax.plot(degrees_array, cv_train_scores, c='C1', lw=3, label='train')
37 ax.set_xlabel('degree')
38 ax.set_ylabel('MSE')
39 ax.set_xticks(degrees_array)
40 fig.savefig('CM3_cv_figure.pdf', format='pdf')
```

We're using several scikit-learn important concepts : transformers, pipelines, grid search

Cross-validation

(3) Choosing the best order for the polynomial approximating a dataset



We're using several scikit-learn important concepts : transformers, pipelines, grid search

- What makes a model **good**?

- Estimating the quality of a model

-  ○ Comparing and selecting models

Comparing and selecting models

Suppose we are given a dataset with p predictors.

We want to estimate a linear model with only a subset of them.

There are mainly **three strategies** for doing this properly:

- **Shrinkage** fits a model with all p predictors but using a modified loss function that drives some parameters to zero automatically. (Razan will talk about this)
- We can use **dimensionality reduction** techniques to project the p predictors to a lower-dimensional subspace. (We will see how to do this in CM4 and TP2)
- In **subset selection** we identify a subset of the predictors that seems the most adequate and then fit a model with them.

Comparing and selecting models

When p is large, testing all 2^p possible models can be very time consuming...

So we prefer to proceed greedily with e.g. **forward stepwise selection**

(1) Let \mathcal{M}_0 denote a model with no predictors (i.e. just the intercept)

(2) For $k = 0, \dots, p - 1$

(a) Consider all $p - k$ models that augment the predictors in \mathcal{M}_k by one extra predictor

(b) Choose the **best** among these $p - k$ models and call it \mathcal{M}_{k+1}

(3) Select single **best** model among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using cross-validated prediction error

Example: Consider the mtcars dataset – we want to predict **mpg**

Description:

The data was extracted from the 1974 _Motor Trend_ US magazine, and comprises fuel consumption and 10 aspects of automobile design and performance for 32 automobiles (1973-74 models).

Format:

A data frame with 32 observations on 11 variables.


[, 1]	mpg	Miles/(US) gallon
[, 2]	cyl	Number of cylinders
[, 3]	disp	Displacement (cu.in.)
[, 4]	hp	Gross horsepower
[, 5]	drat	Rear axle ratio
[, 6]	wt	Weight (lb/1000)
[, 7]	qsec	1/4 mile time
[, 8]	vs	V/S
[, 9]	am	Transmission (0 = automatic, 1 = manual)
[,10]	gear	Number of forward gears
[,11]	carb	Number of carburetors

OLS Regression Results


Dep. Variable:	mpg	R-squared:	0.869
Model:	OLS	Adj. R-squared:	0.807
Method:	Least Squares	F-statistic:	13.93
Date:	Fri, 27 Dec 2024	Prob (F-statistic):	3.79e-07
Time:	15:09:11	Log-Likelihood:	-69.855
No. Observations:	32	AIC:	161.7
Df Residuals:	21	BIC:	177.8
Df Model:	10		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
cyl	-0.1114	1.045	-0.107	0.916	-2.285	2.062
disp	0.0133	0.018	0.747	0.463	-0.024	0.050
hp	-0.0215	0.022	-0.987	0.335	-0.067	0.024
drat	0.7871	1.635	0.481	0.635	-2.614	4.188
wt	-3.7153	1.894	-1.961	0.063	-7.655	0.224
qsec	0.8210	0.731	1.123	0.274	-0.699	2.341
vs	0.3178	2.105	0.151	0.881	-4.059	4.694
am	2.5202	2.057	1.225	0.234	-1.757	6.797
gear	0.6554	1.493	0.439	0.665	-2.450	3.761
carb	-0.1994	0.829	-0.241	0.812	-1.923	1.524
intercept	12.3034	18.718	0.657	0.518	-26.623	51.229


Forward stepwise selection



Forward stepwise selection



Forward stepwise selection



Reminder

We have a TD this Friday at 17h00 !

